

Regression

Jiaren Zhou, B.Eng.

E-Mail: iuk.zhou@gmail.com

Betreuer: Florian Particke M.Sc.

Prof. Dr.-Ing. Jörn Thielecke

Professur für Informationstechnik

Schwerpunkt Navigation und Ortsbestimmung

Telefon: +49 9131 8525-118

Fax: +49 9131 8525-102

E-Mail: joern.thielecke@fau.de

LIKE


Contents

- I. Regression vs. Classification
- II. Linear Regression
- III. Locally Weighted Linear Regression
- IV. Tree-based Regression
- V. Linear Regression vs. Tree Regression
- VI. Summary


I. Regression vs. Classification

"Classification" is the task of predicting a discrete class label.

"Regression" is the task of predicting a continuous quantity.

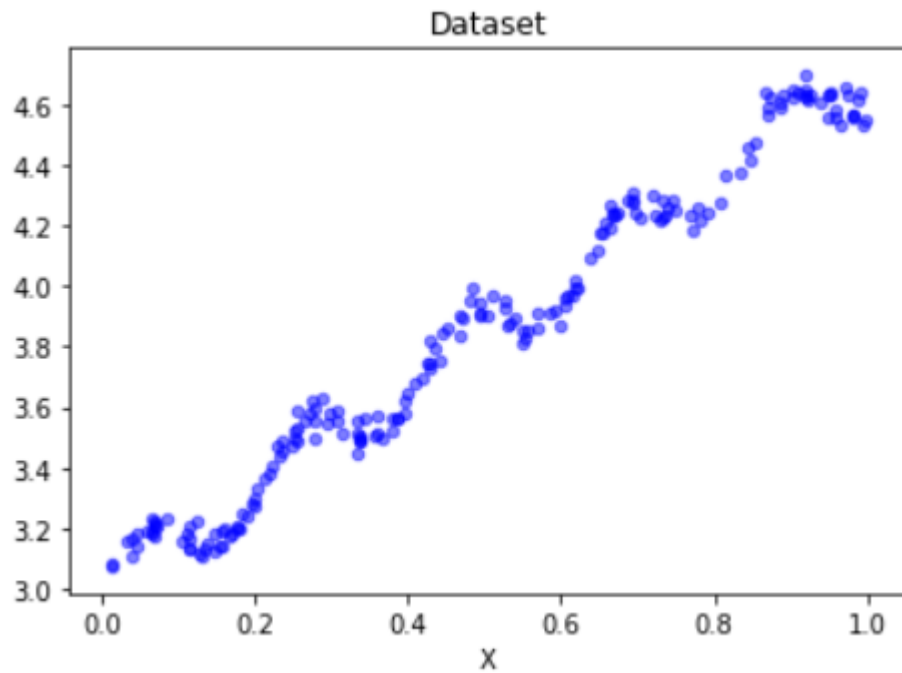


Classification



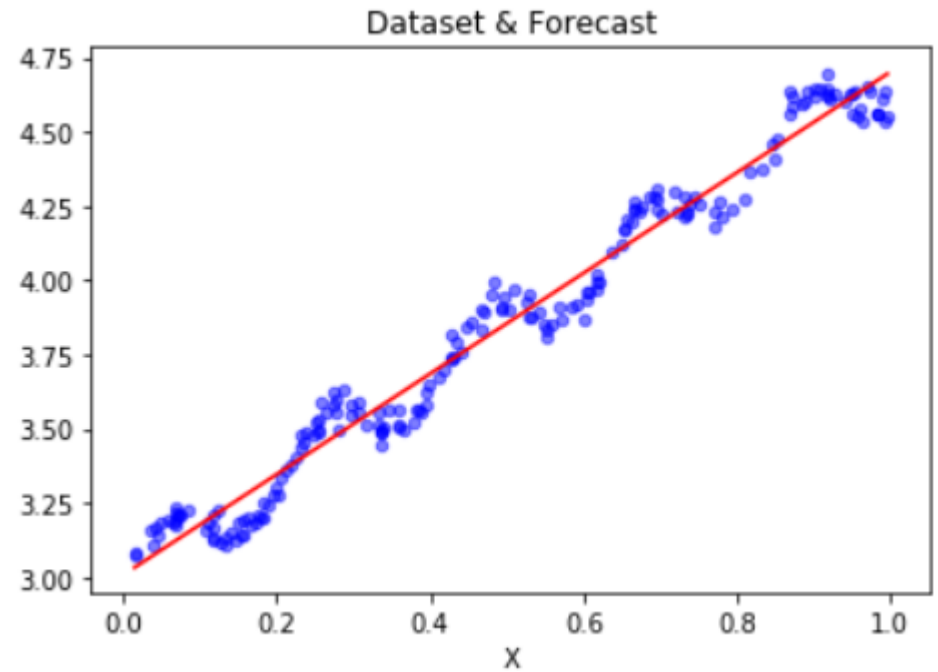
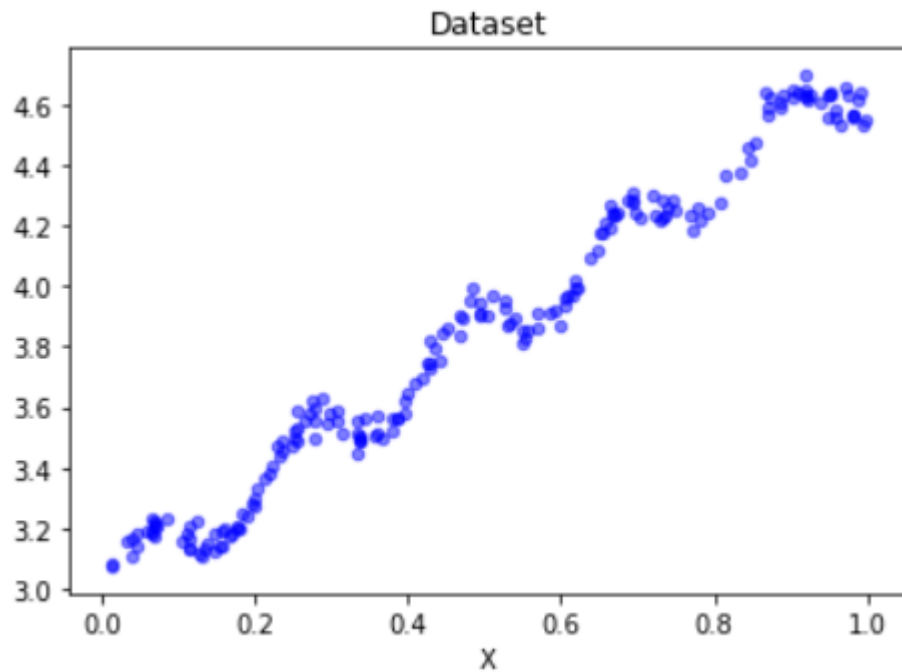
Regression

II. Linear Regression



LIKE

II. Linear Regression



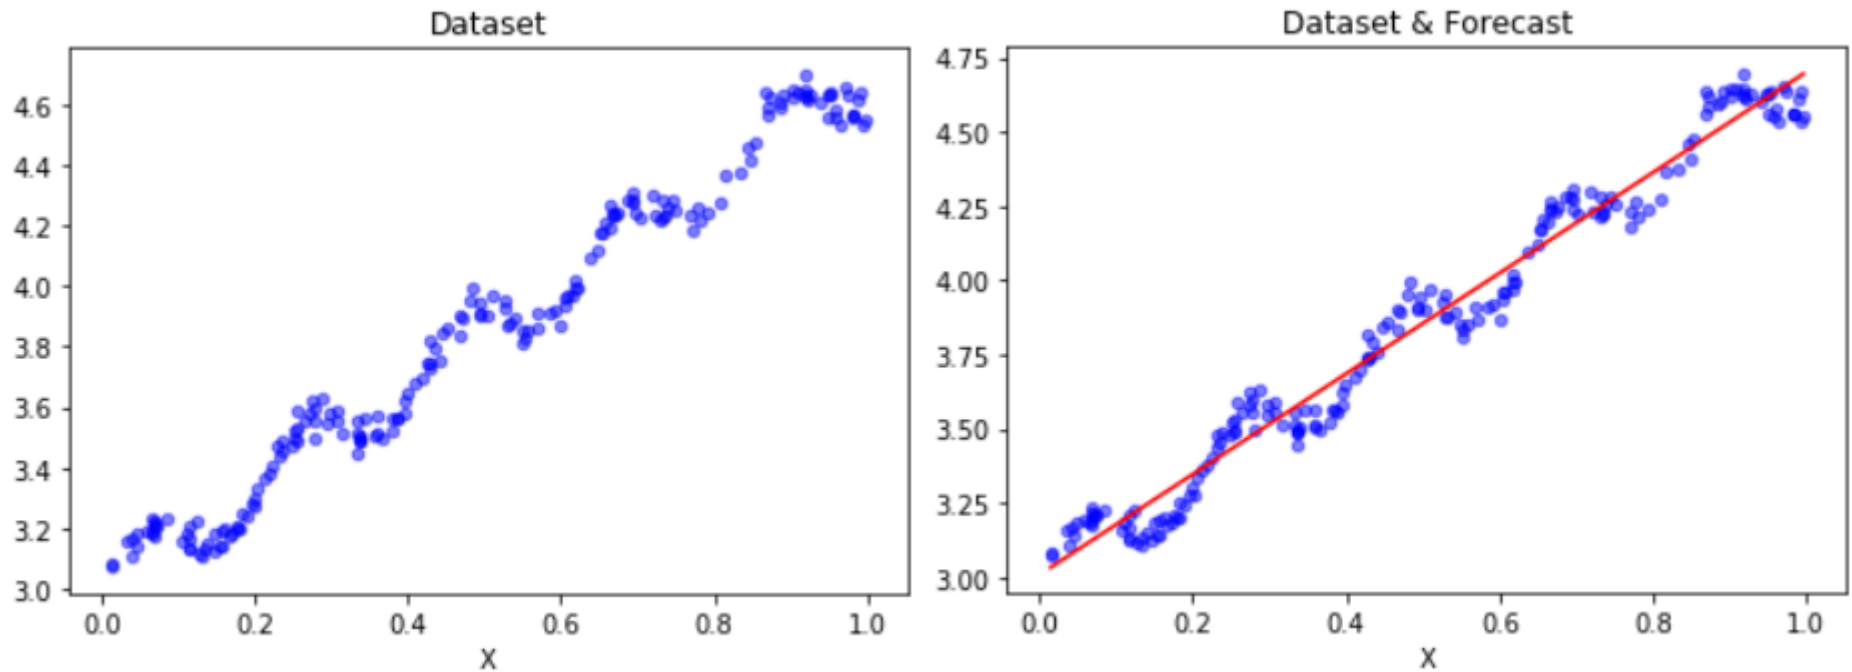
$$y = \omega_1 * x + \omega_0$$

$$\omega_1 : 1.7$$

$$\omega_0 : 3.0$$

LIKE

II. Linear Regression



$$y = \omega_1 * x + \omega_0$$

how to find the best **weight** ?

ω_1 : 1.7

ω_0 : 3.0

LIKE

II. Linear Regression

How to find the best **weight**:

find the least squared error

$$\sum_{i=1}^m (y_i - x_i^T \omega)^2$$

The expression in matrix is

$$(Y - X\omega)^T (Y - X\omega)$$

After the derivation with respect to ω , we get


$$X^T (Y - X\omega)$$

Setting this to zero and solve for ω to get the following equation

$$\hat{\omega} = (X^T X)^{-1} X^T Y$$

LIKE

II. Linear Regression



$$\hat{\omega} = (X^T X)^{-1} X^T Y$$




$$Y = w_1 * X + w_0$$

weights: $\begin{bmatrix} 3.00774324 \\ 1.69532264 \end{bmatrix}$

$w_0 : 3.0$
 $w_1 : 1.7$

LIKE

II. Linear Regression

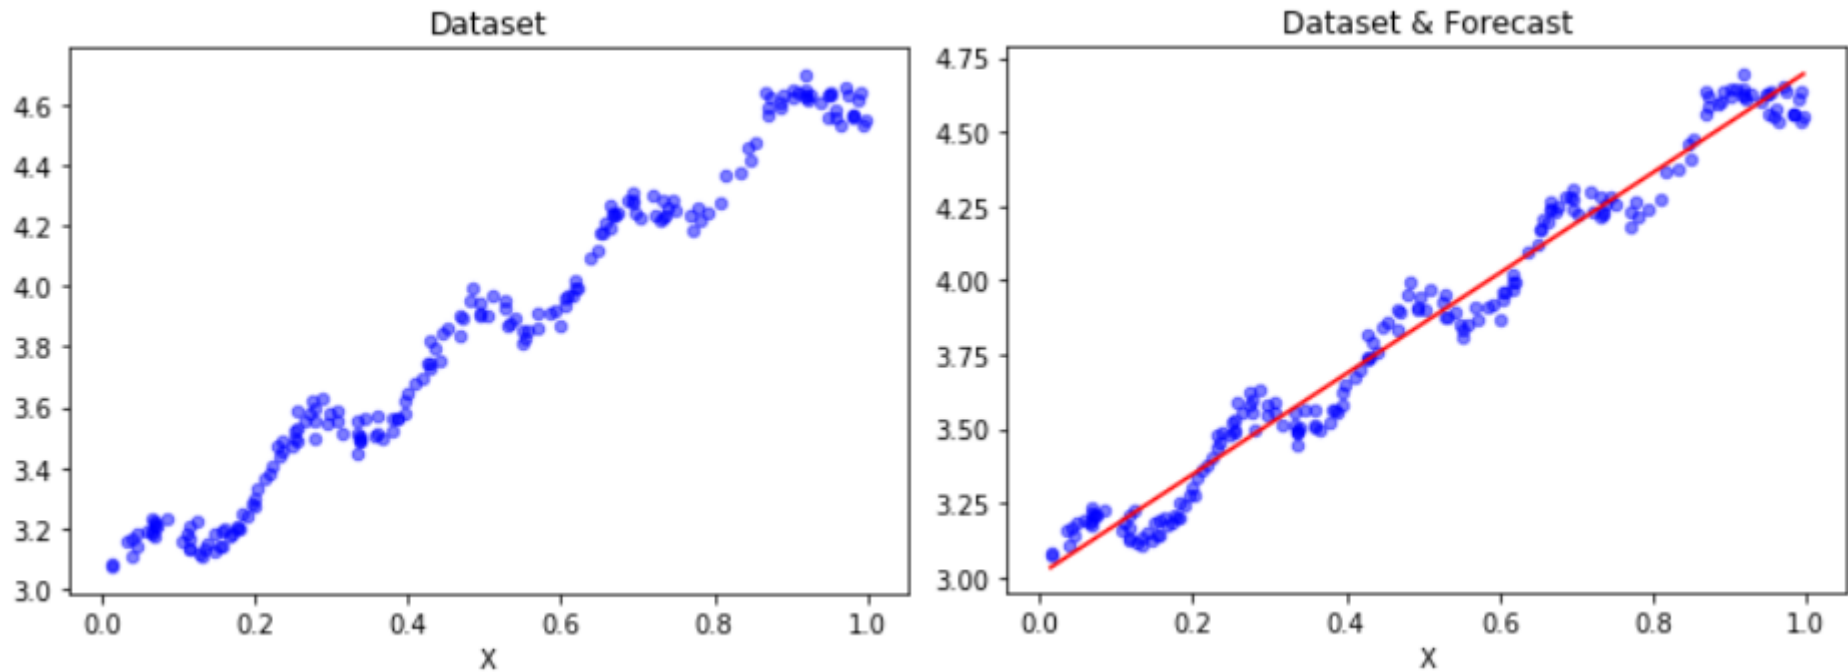


But how far do they differ from each other?

```
numpy.corrcoef(pred, actual):      array([[1.          ,  0.98647356],  
                                     [0.98647356,  1.          ]])
```

LIKE

II. Linear Regression




But how far do they differ from each other?

```
numpy.corrcoef(pred, actual):      array([[1.          ,  0.98647356],  
                                         [0.98647356,  1.          ]])
```

More precise ?

LIKE

II. Linear Regression




But how far do they differ from each other?

```
numpy.corrcoef(pred, actual):      array([[1.          ,  0.98647356],  
                                         [0.98647356,  1.          ]])
```

More precise ?

LIKE

II. Linear Regression




But how far do they differ from each other?

```
numpy.corrcoef(pred, actual):      array([[1.          ,  0.98647356],  
                                     [0.98647356,  1.          ]])
```

More precise ?

LIKE

II. Linear Regression



But how far do they differ from each other?

```
numpy.corrcoef(pred, actual):      array([[1.          ,  0.98647356],  
                                         [0.98647356,  1.          ]])
```

More precise ?

LIKE

III. Locally Weighted Linear Regression

More weights are given to those data points which are close to the data point of interest, then the least-squares regression similar to the linear regression will be carried out.

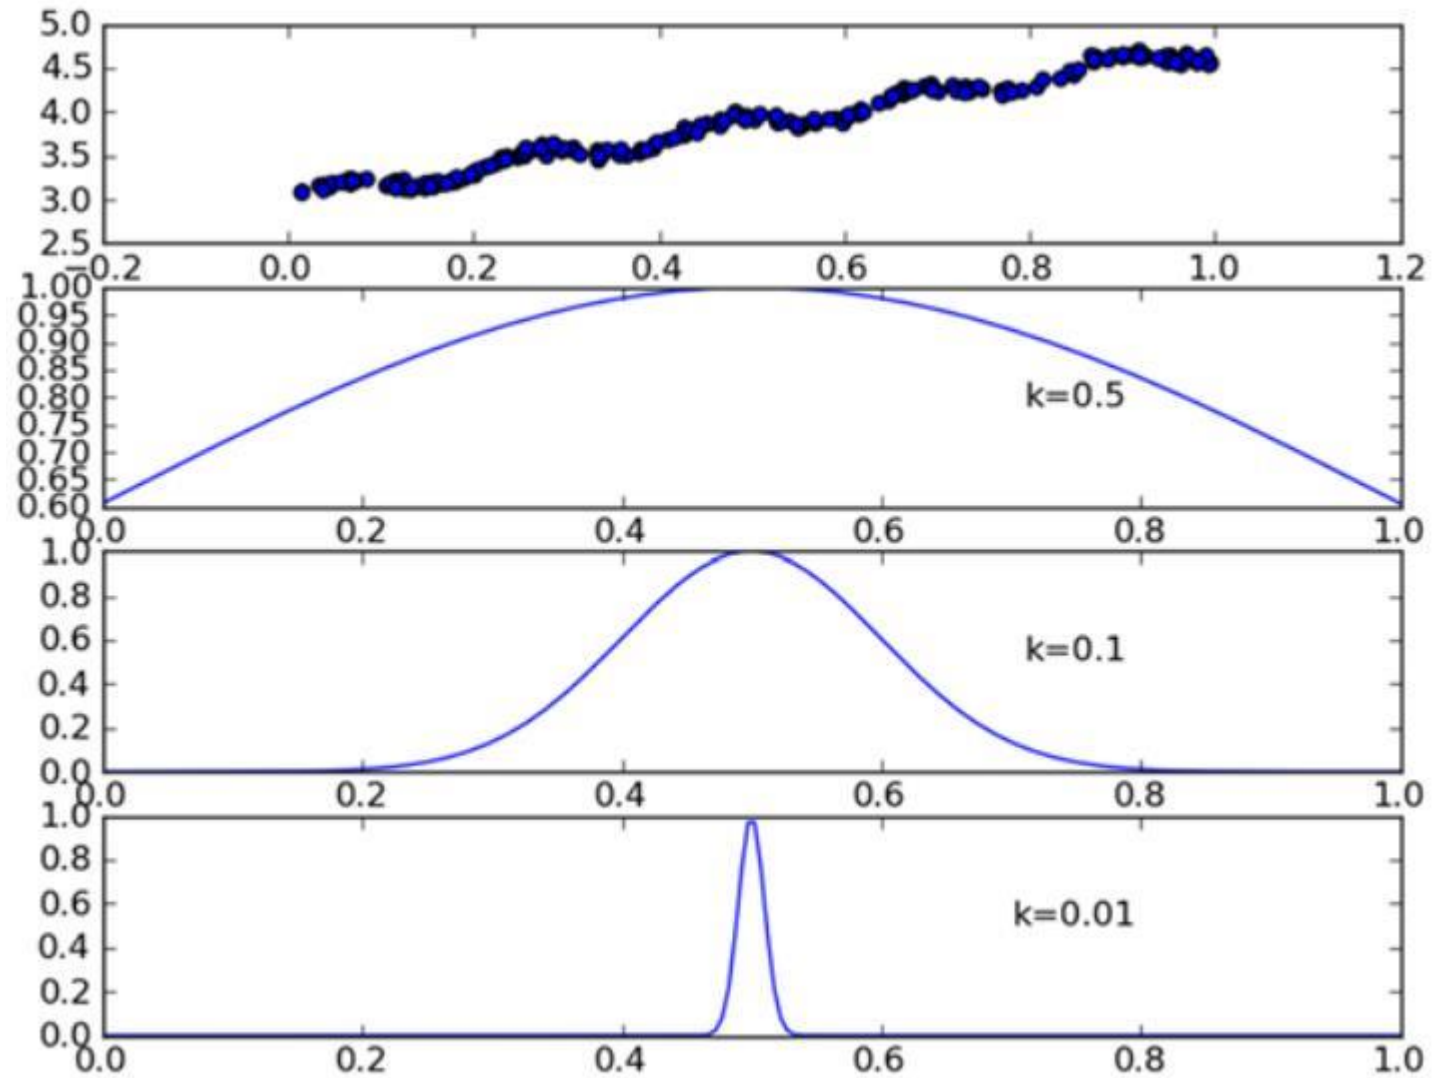
$$\hat{\omega} = (X^T X)^{-1} X^T Y \quad \longrightarrow \quad \hat{\omega} = (X^T W X)^{-1} X^T W Y$$

W is a matrix and will be generated by a **kernel** function, which shall give nearby points more weights than other points. The mostly used kernel is Gaussian and assigns the weights by

$$\omega(i, i) = \exp\left(\frac{|x^i - x|}{-2\kappa^2}\right)$$


III. Locally Weighted Linear Regression

$$\omega(i, i) = \exp\left(\frac{|x^i - x|}{-2\kappa^2}\right)$$



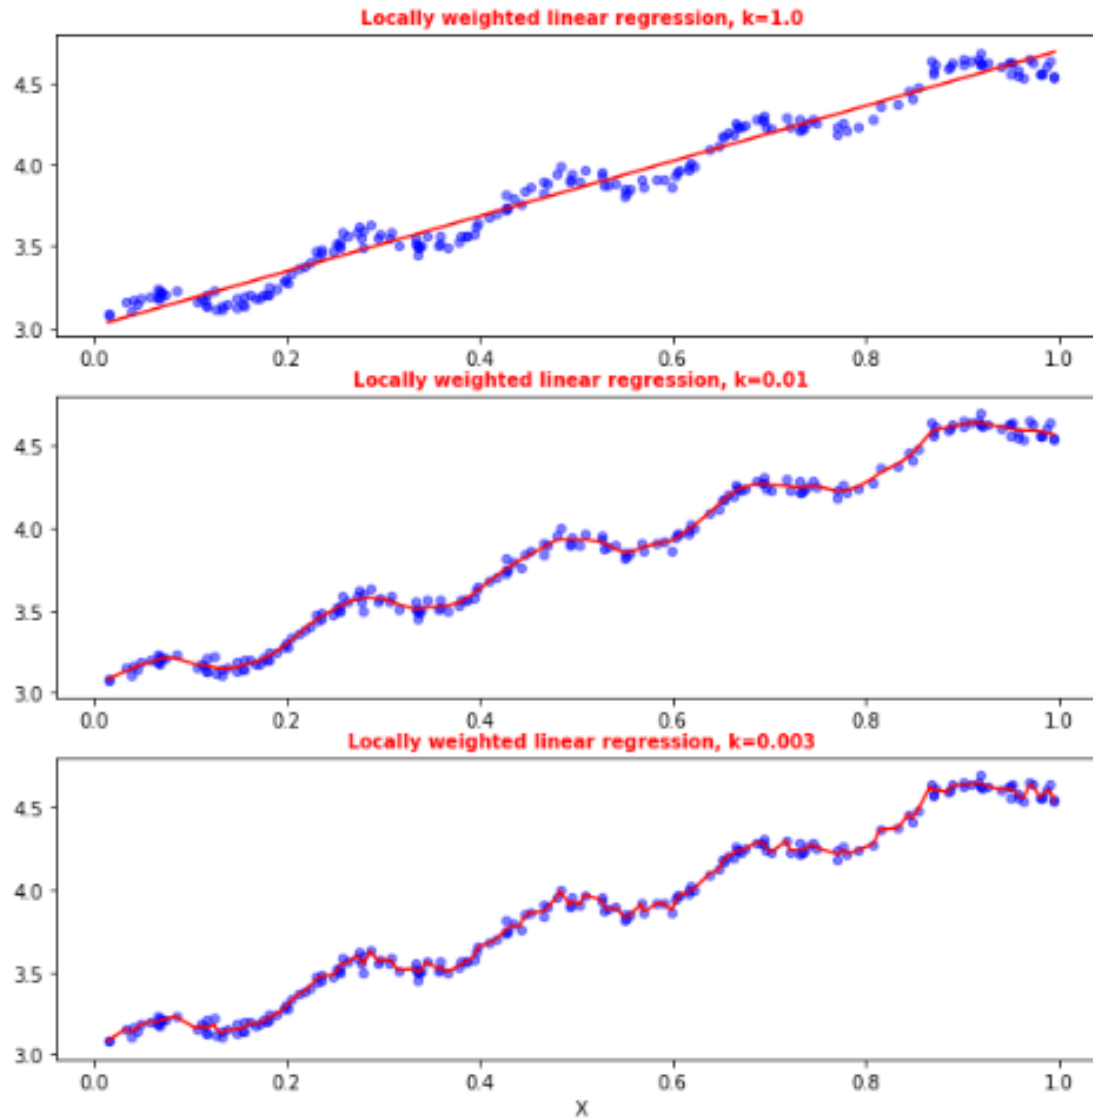
LIKE

III. Locally Weighted Linear Regression



LIKE

III. Locally Weighted Linear Regression



overfitting!

LIKE

Example: predicting the age of an abalone

Testdata: abalone.txt

Feature_X:

```
[[ 1.  0.455  0.365  0.095  0.514  0.2245  0.101  0.15 ]  
[ 1.   0.35  0.265   0.09  0.225  0.0995  0.048  0.07 ]  
[-1.   0.53   0.42  0.135  0.677  0.2565  0.141  0.21 ]  
[ ...   ...   ...   ...   ...   ...   ...   ... ]  
[ 0.   0.33  0.255   0.08  0.205  0.0895  0.039  0.05 ]]
```

Age_Y:

```
[[15.]  
[ 7.]  
[ 9.]  
[10.]  
[ 7.]]
```

Training set and test set are *identical*:

k=0.1, squared error : 56.82523568972884

k=1.0, squared error : 429.8905618700651

k=10, squared error : 549.1181708826451

Example: predicting the age of an abalone

Training set and test set are *identical*:

k=0.1, squared error : 56.82523568972884

k=1.0, squared error : 429.8905618700651

k=10, squared error : 549.1181708826451

Training set and test set are *different*:

k=0.1, squared error : 41317.161723642595


k=1.0, squared error : 573.526144189767

k=10, squared error : 517.5711905387598

Linear regression: squared error : 518.6363153249638

IV: Tree-based Regression

How to deal with the *nonlinearities* in real life?



LIKE


IV: Tree-based Regression

How to deal with the *nonlinearities* in real life?

Tree Regression


IV: Tree-based Regression

Tree structure

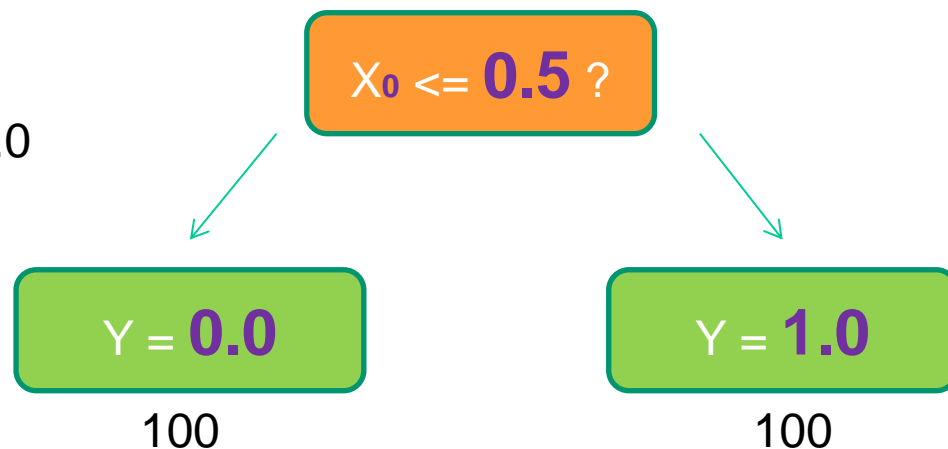


Note:- A is parent node of B and C.

IV: Tree-based Regression




Feature Index: 0
Threshold value: 0.5
Leaf node value: 0.0 or 1.0
Number of Samples: 100




LIKE

IV: Tree-based Regression



How to find the **splitting feature** and **splitting point** ?

Feature Index: 0
Threshold value: 0.5
Leaf node value: 0.0 or 1.0
Number of Samples: 100



LIKE

IV: Tree-based Regression

How to deal with the *nonlinearities* in real life?

Tree regression:

It implements a new algorithm called **CART** (**C**lassification **A**nd **R**egression **T**rees). It is well-known and well-documented tree-building algorithm that makes *binary splits* to handle continuous variables.

By doing this we choose a **feature** and make **values** *greater* than the desired go on the **right** side of the tree and all the other values go on the **left** side.

LIKE

IV: Tree-based Regression

How to make a binary split?

We need to select: splitting feature ' x_j ' and a splitting point ' s ' so that we can divide data into two regions R_1 and R_2 :

$$R_1(j, s) = \{x | x^{(j)} \leq s\}, \quad R_2(j, s) = \{x | x^{(j)} > s\}$$

Then we calculate the average value for each generated region by

$$\hat{c}_1 = \text{ave}(y_i | x_i \in R_1(j, s)), \quad \hat{c}_2 = \text{ave}(y_i | x_i \in R_2(j, s))$$


Goal: find such (\hat{c}_1, \hat{c}_2) which gives the **minimum** of total squared error as follows

$$\min_{j,s} [\min_{\hat{c}_1} \sum_{x_i \in R_1(j,s)} (y_i - \hat{c}_1)^2 + \min_{\hat{c}_2} \sum_{x_i \in R_2(j,s)} (y_i - \hat{c}_2)^2]$$

Example: binary split in the tree regression

Let's split it by the value of the **second** feature,
threshold value: **0.5**

```
matrix([[1., 0., 0., 0.],  
        [0., 1., 0., 0.],  
        [0., 0., 1., 0.],  
        [0., 0., 0., 1.]])
```




Left: $\begin{bmatrix} 1. & 0. & 0. & 0. \\ 0. & 0. & 1. & 0. \\ 0. & 0. & 0. & 1. \end{bmatrix}$



Right: $\begin{bmatrix} 0. & 1. & 0. & 0. \end{bmatrix}$

IV: Tree-based Regression




Stop Conditions:

Mind. error reduction: 1

Mind. data instances: 4

IV: Tree-based Regression




```
print(createTree(data_mat))  
{'splnd': 0,  
'spVal': 0.48813,  
'left': -0.04465028571428572,  
'right': 1.0180967672413792 }
```

Stop Conditions:



Mind. error reduction: 1

Mind. data instances: 4



LIKE


IV: Tree-based Regression



Stop Conditions:


Mind. error reduction: 1

Mind. data instances: 4



LIKE

IV: Tree-based Regression




Stop Conditions:

Mind. error reduction: 1

Mind. data instances: 4


IV: Tree-based Regression



Stop Conditions:

Mind. error reduction: 1

Mind. data instances: 4






IV: Tree-based Regression

Stop Conditions:

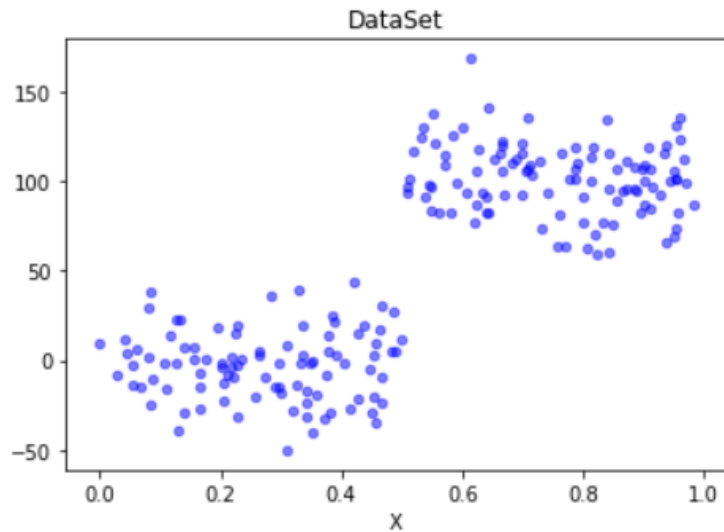
Mind. error reduction: 1

Mind. data instances: 4



LIKE

IV: Tree-based Regression




Stop Conditions:

Mind. error reduction: 1

Mind. data instances: 4

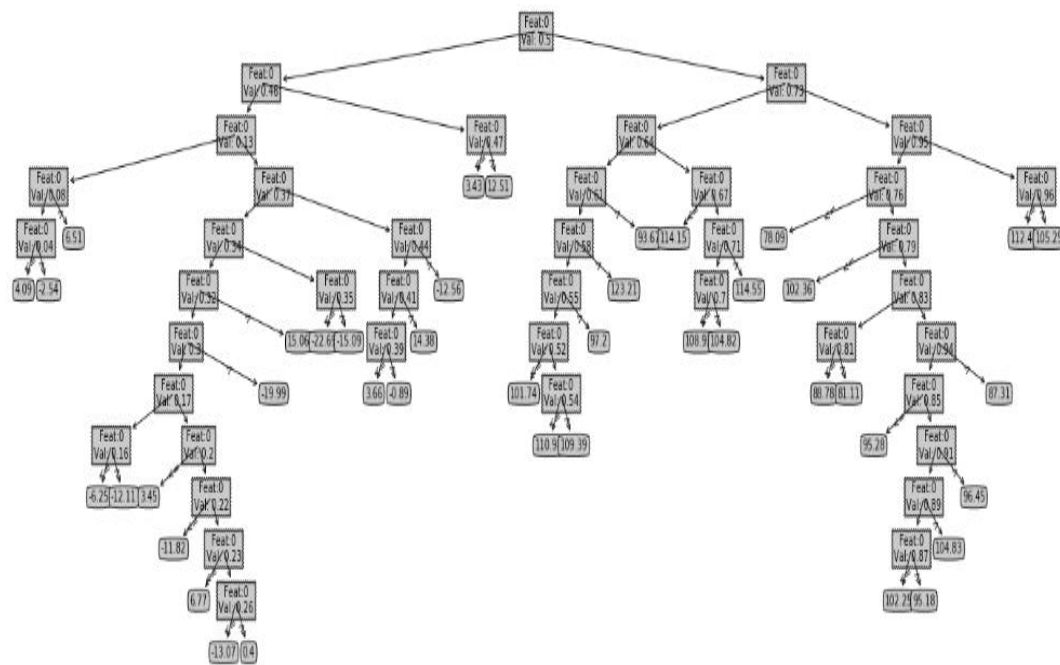
IV: Tree-based Regression





Stop Conditions:

Mind. error reduction: 1

Mind. data instances: 4



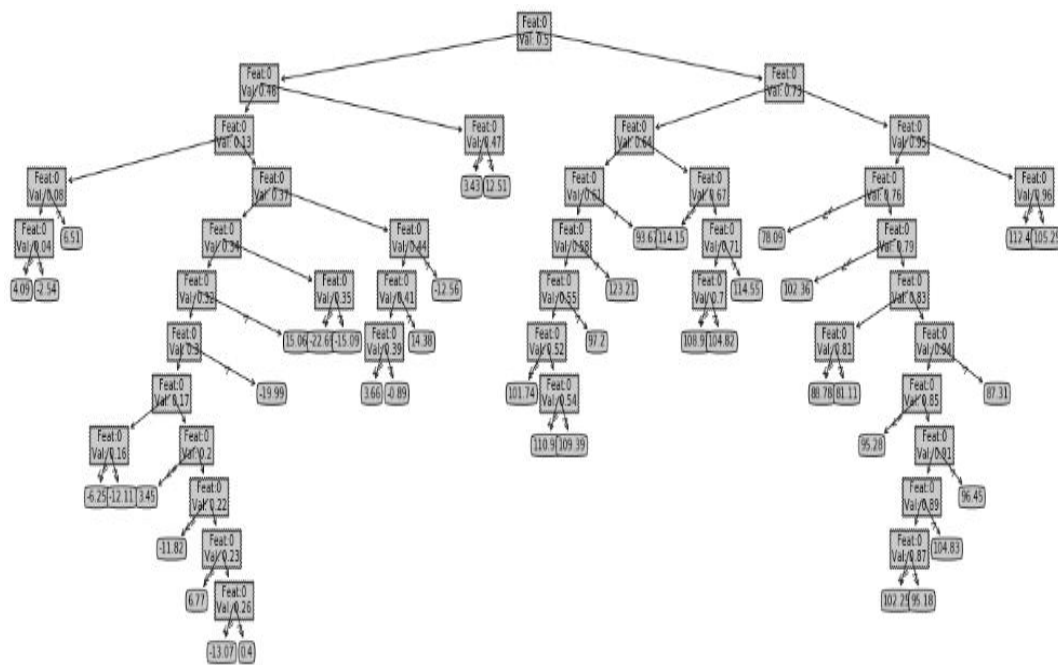
IV: Tree-based Regression



Stop Conditions:



Mind. error reduction: 1

Mind. data instances: 4



LIKE


IV: Tree-based Regression



Stop Conditions:

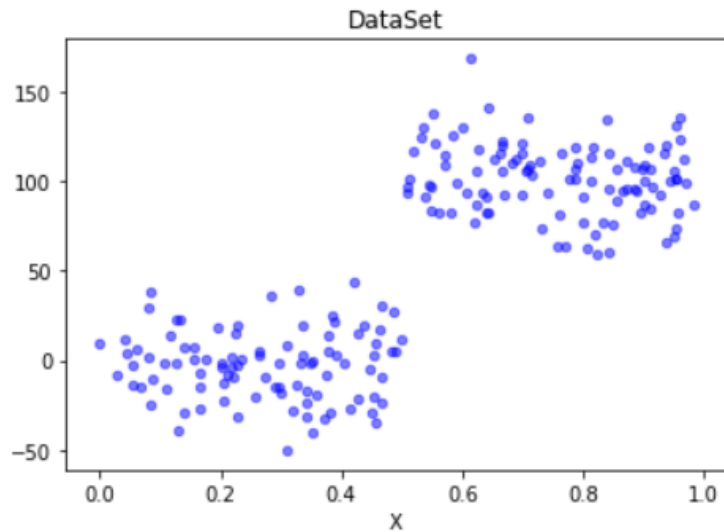
Mind. error reduction: 1

Mind. data instances: 4



LIKE

IV: Tree-based Regression





Stop Conditions:

Mind. error reduction: 4 10000

Mind. data instances: 4


IV: Tree-based Regression



Stop Conditions:



Mind. error reduction: 4 10000

Mind. data instances: 4



LIKE

IV: Tree-based Regression



Stop Conditions:

Mind. error reduction: 4 10000

Mind. data instances: 4

Is there any solution
without user intervention?

→ Postpruning



LIKE


How to implement the postpruning?

- Make the regression tree to a large depth
- Start at the bottom and combine every two leaf nodes(left and right) into a new terminal node, if such error after merge is smaller than the original one

Postpruning

How to implement the postpruning?

- Make the regression tree to a large depth
- Start at the bottom and combine every two leaf nodes(left and right) into a new terminal node, if such error after merge is smaller than the original one




LIKE

Postpruning

How to implement the postpruning?

- Make the regression tree to a large depth
- Start at the bottom and combine every two leaf nodes(left and right) into a new terminal node, if such error after merge is smaller than the original one



Test_data



Calculate Merge-Error
with 45

Calculate Orig.-Error
with 40 and 50


$$\sum_{x_i \in R(2,7)} (y_i - 45)^2 < \sum_{x_i \in R_1(2,7)} (y_i - 40)^2 + \sum_{x_i \in R_2(2,7)} (y_i - 50)^2$$

LIKE

Postpruning

How to implement the postpruning?


- Make the regression tree to a large depth
- Start at the bottom and combine every two leaf nodes(left and right) into a new terminal node, if such error after merge is smaller than the original one



If Merge-Error < Orig.-Error
then **merge**


LIKE

Prepruning vs. Postpruning



Prepruning

Sum of leaf nodes: 200
tree depth: 25





Postpruning

Sum of leaf nodes: 141
tree depth: 23

LIKE

Prepruning vs. Postpruning




Some nodes were pruned off the tree, but it wasn't reduced to two nodes as we had hoped. It turns out that postpruning isn't as effective as prepruning.

We can employ **both** to give the best possible model.

LIKE

IV: Tree-based Regression




Would it be better to model this dataset as

- a bunch of constant values (many leaf nodes) or
- two straight lines (1: from 0.0 to 0.3; 2: from 0.3 to 1.0)

LIKE

IV: Tree-based Regression




Model Tree = Tree + Linear regression

A way to model dataset as a piecewise linear model at each leaf node. Piecewise linear means that the model consists of multiple linear segments

LIKE

IV: Tree-based Regression





Result:

```
{'spInd': 0, 'spVal': 0.285477,  
'left': matrix([[3.46877936], [1.18521743]]),  
'right': matrix([[1.69855694e-03], [1.19647739e+01]])}
```

$$y = 0 + 11.96 * x \quad \text{and} \quad y = 3.47 + 1.20 * x$$

LIKE

IV: Tree-based Regression



Result:

`{'spInd': 0, 'spVal': 0.285477,`

`'left': matrix([[3.46877936], [1.18521743]]),`


`'right': matrix([[1.69855694e-03], [1.19647739e+01]])}`

$$y = 0 + 11.96 * x \quad \text{and} \quad y = 3.47 + 1.20 * x$$


LIKE

V: Linear Regression vs. Tree Regression

Training-data



Testing-data



Evaluation method:


- Correlation coefficients

Regression methods:


- Linear Regression(LR)
- Regression Tree (RT)
- Model Tree (MT)

V: Linear Regression vs. Tree Regression


Linear Regression



Regression Tree



Model Tree



Correlation coefficients:

- LR: 0.94346842356
- RT: 0.96408523182
- MT: 0.97604121913

LIKE

VI: Summary

Regression is the process of predicting a target value, which makes it become one of the most useful tools in statistics.

Linear methods:

- Linear regression: does a good job, when the dataset is simple and linear.
- Locally weighted linear regression: the forecast would be more precise. Overfitting problem should be avoided.

Non-linear methods:

- Regression tree: breaks up the predicted value into piecewise constant segments.
- Model Tree: implements the linear regression equations at each leaf node.
- Pre- and Postpruning: can effectively reduce the complexity of tree and help avoid the overfitting problem.

Reference

- ❖ Harrington, Peter. 2012. *Machine Learning in Action*. Shelter Island (N.Y.): Manning Publications Co.
- ❖ Li Hang. 2012. *Statistic Learning Methods*. Tsinghua University Publications Co.
- ❖ Analytics Vidhya: A Complete Tutorial on Tree Based Modeling from Scratch. APRIL 12, 2016
<https://www.analyticsvidhya.com/blog/2016/04/complete-tutorial-tree-based-modeling-scratch-in-python>